

MODULE 4 - HANDOFF

Agent Role Decision Matrix

Claude / Codex / ChatGPT / Cursor - which one for which work type.

Version: v1 - 2026-05-09

Author: Chase Whited / Whited Consulting

Stable URL: raw.githubusercontent.com/cwhited26/aipocketagency-brain

Agent Role Decision Matrix

When work needs an agent, which one do you reach for? This matrix is the working answer. Adjust the columns to your stack; the rows are the work types that recur.

Work type	Claude (Sonnet/Opus)	Codex / Claude Code	ChatGPT	Cursor
Architecture / system design	Primary. Long-context reasoning over the whole brain.	Secondary. Use only after a plan is locked.	Tertiary. Useful for second opinions.	Avoid. The IDE-tight loop is wrong shape for design.
Spec / decision doc writing	Primary. Decision Log entries, ADRs, lane plans.	Avoid. Codex is for execution, not synthesis.	Secondary. Long-form writing where length is the point.	Avoid.
Code execution (commits, migrations, tests)	Secondary. Capable but slower than Codex/Claude Code.	Primary. Built for the execute-and-commit loop.	Avoid. Not built for filesystem writes.	Secondary . Tight IDE feedback, smaller batches.
Pair-programming / fast feedback in IDE	Avoid. The brain reads slow this loop down.	Secondary.	Avoid.	Primary. The native IDE integration is its job.
Customer-facing copy (email, sales, post)	Secondary. Capable but ChatGPT is faster on this surface.	Avoid.	Primary. Long-form, voice-aware writing.	Avoid.
Research / summarization across many sources	Primary. Long context handles the source pile.	Avoid.	Secondary. Faster on small piles.	Avoid.
Code review	Primary. Strong at finding logic bugs and design drift.	Secondary. Better at "does it run" than "is it right."	Tertiary.	Secondary . Inline review in the IDE works for small diffs.

Work type	Claude (Sonnet/Opus)	Codex / Claude Code	ChatGPT	Cursor
Multi-step orchestration overnight	Secondary.	Secondary.	Avoid.	Avoid.
Schema design / data modeling	Primary. Hold the whole schema in context.	Secondary. Migration writing.	Tertiary.	Avoid.
Debug a stack trace / production incident	Primary. Reasoning over symptoms + repo.	Secondary. After diagnosis, executes the fix.	Tertiary.	Secondary . Inline debugger pairing.

Notes on the choices

- **Claude as planner.** Long context windows let Claude hold an entire repo plus the brain and reason across files. Use it to decide; use Codex to do.
- **Codex / Claude Code as executor.** Built for the commit-and-verify loop. Hand it a plan; do not ask it to design.
- **ChatGPT for long-form writing.** Sales pages, customer emails, narrative summaries. Voice-aware when seeded with feedback.md tone entries.
- **Cursor for IDE-tight work.** Pair programming, refactors that need fast iteration, small diffs you can review in-line.
- **Manus / Dispatch for orchestration.** Long-running batch work where you want progress while you sleep. Treat its outputs as untrusted until reviewed.

How to use the matrix

Pick the work type. Pick the agent in the "Primary" column. If primary is unavailable (rate-limited, offline, between releases), drop to secondary. Avoid means real cost - either the model isn't strong on this surface, or the tooling fights the loop.

When the matrix says one thing and your gut says another, follow your gut. The matrix is a starting point. Edit your copy whenever a different agent wins on a real task you ran.